

# dataScience-with-answers

September 22, 2017

## 1 Python for Data Analysis

**Research Computing Services** Website: [rcs.bu.edu](http://rcs.bu.edu) Tutorial materials: [http://rcs.bu.edu/examples/python/data\\_analysis](http://rcs.bu.edu/examples/python/data_analysis)

```
In [1]: #Import Python Libraries
import numpy as np
import scipy as sp
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: #Read csv file
df = pd.read_csv("http://rcs.bu.edu/examples/python/data_analysis/Salaries.csv")
```

```
In [3]: #Display a few first records
df.head()
```

```
Out[3]:
```

	rank	discipline	phd	service	sex	salary
0	Prof	B	56	49	Male	186960
1	Prof	A	12	6	Male	93000
2	Prof	A	23	20	Male	110515
3	Prof	A	40	31	Male	131205
4	Prof	B	20	18	Male	104800

---

*Excercise*

```
In [4]: #Display first 10 records
# <your code goes here>
df.head(10)
```

```
Out[4]:
```

	rank	discipline	phd	service	sex	salary
0	Prof	B	56	49	Male	186960
1	Prof	A	12	6	Male	93000
2	Prof	A	23	20	Male	110515
3	Prof	A	40	31	Male	131205

4	Prof	B	20	18	Male	104800
5	Prof	A	20	20	Male	122400
6	AssocProf	A	20	17	Male	81285
7	Prof	A	18	18	Male	126300
8	Prof	A	29	19	Male	94350
9	Prof	A	51	51	Male	57800

```
In [5]: #Display first 20 records
# <your code goes here>
df.head(20)
```

```
Out[5]:
```

	rank	discipline	phd	service	sex	salary
0	Prof	B	56	49	Male	186960
1	Prof	A	12	6	Male	93000
2	Prof	A	23	20	Male	110515
3	Prof	A	40	31	Male	131205
4	Prof	B	20	18	Male	104800
5	Prof	A	20	20	Male	122400
6	AssocProf	A	20	17	Male	81285
7	Prof	A	18	18	Male	126300
8	Prof	A	29	19	Male	94350
9	Prof	A	51	51	Male	57800
10	Prof	B	39	33	Male	128250
11	Prof	B	23	23	Male	134778
12	AsstProf	B	1	0	Male	88000
13	Prof	B	35	33	Male	162200
14	Prof	B	25	19	Male	153750
15	Prof	B	17	3	Male	150480
16	AsstProf	B	8	3	Male	75044
17	AsstProf	B	4	0	Male	92000
18	Prof	A	19	7	Male	107300
19	Prof	A	29	27	Male	150500

```
In [6]: #Display the last 5 records
# <your code goes here>
df.tail()
```

```
Out[6]:
```

	rank	discipline	phd	service	sex	salary
73	Prof	B	18	10	Female	105450
74	AssocProf	B	19	6	Female	104542
75	Prof	B	17	17	Female	124312
76	Prof	A	28	14	Female	109954
77	Prof	A	23	15	Female	109646

---

```
In [7]: #Identify the type of df object
type(df)
```

```

Out[7]: pandas.core.frame.DataFrame

In [8]: #Check the type of a column "salary"
        df['salary'].dtype

Out[8]: dtype('int64')

In [9]: #List the types of all columns
        df.dtypes

Out[9]: rank          object
        discipline    object
        phd           int64
        service       int64
        sex           object
        salary        int64
        dtype: object

In [10]: #List the column names
         df.columns

Out[10]: Index(['rank', 'discipline', 'phd', 'service', 'sex', 'salary'], dtype='object')

In [11]: #List the row labels and the column names
         df.axes

Out[11]: [RangeIndex(start=0, stop=78, step=1),
         Index(['rank', 'discipline', 'phd', 'service', 'sex', 'salary'], dtype='object')]

In [12]: #Number of dimensions
         df.ndim

Out[12]: 2

In [13]: #Total number of elements in the Data Frame
         df.size

Out[13]: 468

In [14]: #Number of rows and columns
         df.shape

Out[14]: (78, 6)

In [15]: #Output basic statistics for the numeric columns
         df.describe()

Out[15]:
```

	phd	service	salary
count	78.000000	78.000000	78.000000
mean	19.705128	15.051282	108023.782051
std	12.498425	12.139768	28293.661022

min	1.000000	0.000000	57800.000000
25%	10.250000	5.250000	88612.500000
50%	18.500000	14.500000	104671.000000
75%	27.750000	20.750000	126774.750000
max	56.000000	51.000000	186960.000000

```
In [16]: #Calculate mean for all numeric columns
df.mean()
```

```
Out[16]: phd          19.705128
service      15.051282
salary      108023.782051
dtype: float64
```

---

### Excercise

```
In [17]: #Calculate the standard deviation (std() method) for all numeric columns
# <your code goes here>
df.std()
```

```
Out[17]: phd          12.498425
service      12.139768
salary      28293.661022
dtype: float64
```

```
In [18]: #Calculate average of the columns in the first 50 rows
# <your code goes here>
df.head(50).mean()
```

```
Out[18]: phd          21.52
service      17.60
salary      113789.14
dtype: float64
```

---

## 1.0.1 Data slicing and grouping

```
In [19]: df_sex = df.groupby('sex')
```

```
In [20]: #Extract a column by name (method 1)
df['sex'].head()
```

```
Out[20]: 0    Male
1    Male
2    Male
3    Male
4    Male
Name: sex, dtype: object
```

```
In [21]: #Extract a column name (method 2)
df.sex.head()
```

```
Out[21]: 0    Male
         1    Male
         2    Male
         3    Male
         4    Male
         Name: sex, dtype: object
```

---

### Excercise

```
In [22]: #Calculate the basic statistics for the salary column (used describe() method)
         # <your code goes here>
df['salary'].describe()
```

```
Out[22]: count      78.000000
         mean    108023.782051
         std     28293.661022
         min     57800.000000
         25%     88612.500000
         50%    104671.000000
         75%    126774.750000
         max     186960.000000
         Name: salary, dtype: float64
```

```
In [23]: #Calculate how many values in the salary column (use count() method)
         # <your code goes here>
df['salary'].count()
```

```
Out[23]: 78
```

```
In [24]: #Calculate the average salary
df['salary'].mean()
```

```
Out[24]: 108023.78205128205
```

---

```
In [25]: #Group data using rank
df_rank = df.groupby('rank')
```

```
In [26]: #Calculate mean of all numeric columns for the grouped object
df_rank.mean()
```

```
Out[26]:
```

	phd	service	salary
rank			
AssocProf	15.076923	11.307692	91786.230769
AsstProf	5.052632	2.210526	81362.789474
Prof	27.065217	21.413043	123624.804348

```
In [27]: #Calculate the mean salary for men and women. The following produce Pandas Series (sing
df.groupby('sex')['salary'].mean()
```

```
Out[27]: sex
Female    101002.410256
Male      115045.153846
Name: salary, dtype: float64
```

```
In [28]: # If we use double brackets Pandas will produce a DataFrame
df.groupby('sex')[['salary']].mean()
```

```
Out[28]:          salary
sex
Female    101002.410256
Male      115045.153846
```

```
In [29]: # Group using 2 variables - sex and rank:
df.groupby(['sex', 'rank'], sort=False)[['salary']].mean()
```

```
Out[29]:          salary
sex    rank
Male   Prof      124690.142857
        AssocProf 102697.666667
        AsstProf   85918.000000
Female Prof      121967.611111
        AssocProf   88512.800000
        AsstProf   78049.909091
```

---

### Excercise

```
In [30]: # Group data by the discipline and find the average salary for each group
df.groupby('discipline')['salary'].mean()
```

```
Out[30]: discipline
A      98331.111111
B     116331.785714
Name: salary, dtype: float64
```

---

## 1.0.2 Filtering

```
In [31]: #Select observation with the value in the salary column > 120K
df_sub = df[ df['salary'] > 120000]
df_sub.head()
```

```
Out[31]:
```

	rank	discipline	phd	service	sex	salary
0	Prof	B	56	49	Male	186960
3	Prof	A	40	31	Male	131205
5	Prof	A	20	20	Male	122400
7	Prof	A	18	18	Male	126300
10	Prof	B	39	33	Male	128250

```
In [32]: #Select data for female professors
df_w = df[ df['sex'] == 'Female']
df_w.head()
```

```
Out[32]:
```

	rank	discipline	phd	service	sex	salary
39	Prof	B	18	18	Female	129000
40	Prof	A	39	36	Female	137000
41	AssocProf	A	13	8	Female	74830
42	AsstProf	B	4	2	Female	80225
43	AsstProf	B	5	0	Female	77000

---

### Excercise

```
In [33]: # Using filtering, find the mean value of the salary for the discipline A
df[df['discipline'] == 'A']['salary'].mean()
```

```
Out[33]: 98331.1111111111109
```

```
In [34]: # Challenge:
# Extract (filter) only observations with high salary ( > 100K) and find how many female
df[df['salary'] > 120000].groupby('sex')['salary'].count()
```

```
Out[34]: sex
Female      9
Male       16
Name: salary, dtype: int64
```

---

### 1.0.3 More on slicing the dataset

```
In [35]: #Select column salary
df1 = df['salary']
```

```
In [36]: #Check data type of the result
type(df1)
```

```
Out[36]: pandas.core.series.Series
```

```
In [37]: #Look at the first few elements of the output
df1.head()
```

```
Out[37]: 0    186960
         1     93000
         2    110515
         3    131205
         4    104800
         Name: salary, dtype: int64
```

```
In [38]: #Select column salary and make the output to be a data frame
         df2 = df[['salary']]
```

```
In [39]: #Check the type
         type(df2)
```

```
Out[39]: pandas.core.frame.DataFrame
```

```
In [40]: #Select a subset of rows (based on their position):
         # Note 1: The location of the first row is 0
         # Note 2: The last value in the range is not included
         df[0:10]
```

```
Out[40]:
```

	rank	discipline	phd	service	sex	salary
0	Prof	B	56	49	Male	186960
1	Prof	A	12	6	Male	93000
2	Prof	A	23	20	Male	110515
3	Prof	A	40	31	Male	131205
4	Prof	B	20	18	Male	104800
5	Prof	A	20	20	Male	122400
6	AssocProf	A	20	17	Male	81285
7	Prof	A	18	18	Male	126300
8	Prof	A	29	19	Male	94350
9	Prof	A	51	51	Male	57800

```
In [41]: #If we want to select both rows and columns we can use method .loc
         df.loc[10:20,['rank', 'sex', 'salary']]
```

```
Out[41]:
```

	rank	sex	salary
10	Prof	Male	128250
11	Prof	Male	134778
12	AsstProf	Male	88000
13	Prof	Male	162200
14	Prof	Male	153750
15	Prof	Male	150480
16	AsstProf	Male	75044
17	AsstProf	Male	92000
18	Prof	Male	107300
19	Prof	Male	150500
20	AsstProf	Male	92000

```
In [42]: #Let's see what we get for our df_sub data frame
         # Method .loc subset the data frame based on the labels:
         df_sub.loc[10:20,['rank', 'sex', 'salary']]
```



```
Out[42]:
```

	rank	sex	salary
10	Prof	Male	128250
11	Prof	Male	134778
13	Prof	Male	162200
14	Prof	Male	153750
15	Prof	Male	150480
19	Prof	Male	150500

```
In [43]: # Unlike method .loc, method iloc selects rows (and columns) by position:
df_sub.iloc[10:20, [0,3,4,5]]
```

```
Out[43]:
```

	rank	service	sex	salary
26	Prof	19	Male	148750
27	Prof	43	Male	155865
29	Prof	20	Male	123683
31	Prof	21	Male	155750
35	Prof	23	Male	126933
36	Prof	45	Male	146856
39	Prof	18	Female	129000
40	Prof	36	Female	137000
44	Prof	19	Female	151768
45	Prof	25	Female	140096

#### 1.0.4 Sorting the Data

```
In [44]: #Sort the data frame by yrs.service and create a new data frame
df_sorted = df.sort_values(by = 'service')
df_sorted.head()
```

```
Out[44]:
```

	rank	discipline	phd	service	sex	salary
55	AsstProf	A	2	0	Female	72500
23	AsstProf	A	2	0	Male	85000
43	AsstProf	B	5	0	Female	77000
17	AsstProf	B	4	0	Male	92000
12	AsstProf	B	1	0	Male	88000

```
In [45]: #Sort the data frame by yrs.service and overwrite the original dataset
df.sort_values(by = 'service', ascending = False, inplace = True)
df.head()
```

```
Out[45]:
```

	rank	discipline	phd	service	sex	salary
9	Prof	A	51	51	Male	57800
0	Prof	B	56	49	Male	186960
36	Prof	B	45	45	Male	146856
27	Prof	A	45	43	Male	155865
40	Prof	A	39	36	Female	137000

```
In [46]: # Restore the original order (by sorting using index)
df.sort_index(axis=0, ascending = True, inplace = True)
df.head()
```

```
Out[46]:
```

	rank	discipline	phd	service	sex	salary
0	Prof	B	56	49	Male	186960
1	Prof	A	12	6	Male	93000
2	Prof	A	23	20	Male	110515
3	Prof	A	40	31	Male	131205
4	Prof	B	20	18	Male	104800

Excercise

```
In [47]: # Sort data frame by the salary (in descending order) and display the first few records
df.sort_values(by='salary', ascending=False).head()
```

```
Out[47]:
```

	rank	discipline	phd	service	sex	salary
0	Prof	B	56	49	Male	186960
13	Prof	B	35	33	Male	162200
72	Prof	B	24	15	Female	161101
27	Prof	A	45	43	Male	155865
31	Prof	B	22	21	Male	155750

---

```
In [48]: #Sort the data frame using 2 or more columns:
df_sorted = df.sort_values(by = ['service', 'salary'], ascending = [True,False])
df_sorted.head(10)
```

```
Out[48]:
```

	rank	discipline	phd	service	sex	salary
52	Prof	A	12	0	Female	105000
17	AsstProf	B	4	0	Male	92000
12	AsstProf	B	1	0	Male	88000
23	AsstProf	A	2	0	Male	85000
43	AsstProf	B	5	0	Female	77000
55	AsstProf	A	2	0	Female	72500
57	AsstProf	A	3	1	Female	72500
28	AsstProf	B	7	2	Male	91300
42	AsstProf	B	4	2	Female	80225
68	AsstProf	A	4	2	Female	77500

### 1.0.5 Missing Values

```
In [49]: # Read a dataset with missing values
flights = pd.read_csv("http://rcs.bu.edu/examples/python/data_analysis/flights.csv")
flights.head()
```

```
Out[49]:
```

	year	month	day	dep_time	dep_delay	arr_time	arr_delay	carrier	tailnum	\
0	2013	1	1	517.0	2.0	830.0	11.0	UA	N14228	
1	2013	1	1	533.0	4.0	850.0	20.0	UA	N24211	
2	2013	1	1	542.0	2.0	923.0	33.0	AA	N619AA	
3	2013	1	1	554.0	-6.0	812.0	-25.0	DL	N668DN	
4	2013	1	1	554.0	-4.0	740.0	12.0	UA	N39463	

	flight	origin	dest	air_time	distance	hour	minute
0	1545	EWR	IAH	227.0	1400	5.0	17.0
1	1714	LGA	IAH	227.0	1416	5.0	33.0
2	1141	JFK	MIA	160.0	1089	5.0	42.0
3	461	LGA	ATL	116.0	762	5.0	54.0
4	1696	EWR	ORD	150.0	719	5.0	54.0

```
In [50]: # Select the rows that have at least one missing value
flights[flights.isnull().any(axis=1)].head()
```

```
Out [50]:
```

	year	month	day	dep_time	dep_delay	arr_time	arr_delay	carrier	\
330	2013	1	1	1807.0	29.0	2251.0	NaN	UA	
403	2013	1	1	NaN	NaN	NaN	NaN	AA	
404	2013	1	1	NaN	NaN	NaN	NaN	AA	
855	2013	1	2	2145.0	16.0	NaN	NaN	UA	
858	2013	1	2	NaN	NaN	NaN	NaN	AA	

	tailnum	flight	origin	dest	air_time	distance	hour	minute
330	N31412	1228	EWR	SAN	NaN	2425	18.0	7.0
403	N3EHAA	791	LGA	DFW	NaN	1389	NaN	NaN
404	N3EVAA	1925	LGA	MIA	NaN	1096	NaN	NaN
855	N12221	1299	EWR	RSW	NaN	1068	21.0	45.0
858	NaN	133	JFK	LAX	NaN	2475	NaN	NaN

```
In [51]: # Filter all the rows where arr_delay value is missing:
flights1 = flights[ flights['arr_delay'].notnull( )]
flights1.head()
```

```
Out [51]:
```

	year	month	day	dep_time	dep_delay	arr_time	arr_delay	carrier	tailnum	\
0	2013	1	1	517.0	2.0	830.0	11.0	UA	N14228	
1	2013	1	1	533.0	4.0	850.0	20.0	UA	N24211	
2	2013	1	1	542.0	2.0	923.0	33.0	AA	N619AA	
3	2013	1	1	554.0	-6.0	812.0	-25.0	DL	N668DN	
4	2013	1	1	554.0	-4.0	740.0	12.0	UA	N39463	

	flight	origin	dest	air_time	distance	hour	minute
0	1545	EWR	IAH	227.0	1400	5.0	17.0
1	1714	LGA	IAH	227.0	1416	5.0	33.0
2	1141	JFK	MIA	160.0	1089	5.0	42.0
3	461	LGA	ATL	116.0	762	5.0	54.0
4	1696	EWR	ORD	150.0	719	5.0	54.0

```
In [52]: # Remove all the observations with missing values
flights2 = flights.dropna()
```

```
In [53]: # Fill missing values with zeros
nomiss = flights['dep_delay'].fillna(0)
nomiss.isnull().any()
```

Out [53]: False

---

### Excercise

```
In [54]: # Count how many missing data are in dep_delay and arr_delay columns
         flights[['dep_delay', 'arr_delay']].isnull().sum()
```

```
Out [54]: dep_delay    2336
         arr_delay    2827
         dtype: int64
```

---

## 1.0.6 Common Aggregation Functions:

Function	Description
min	minimum
max	maximum
count	number of non-null observations
sum	sum of values
mean	arithmetic mean of values
median	median
mad	mean absolute deviation
mode	mode
prod	product of values
std	standard deviation
var	unbiased variance

```
In [55]: # Find the number of non-missing values in each column
         flights.count()
```

```
Out [55]: year          160754
         month          160754
         day            160754
         dep_time       158418
         dep_delay      158418
         arr_time       158275
         arr_delay      157927
         carrier        160754
         tailnum        159321
         flight         160754
         origin         160754
         dest           160754
         air_time       157927
         distance       160754
```

```
hour          158418
minute        158418
dtype: int64
```

```
In [56]: # Find mean value for all the columns in the dataset
flights.min()
```

```
Out[56]: year          2013
month            1
day              1
dep_time         1
dep_delay       -33
arr_time         1
arr_delay       -75
carrier          AA
flight           1
origin           EWR
dest             ANC
air_time         21
distance         17
hour             0
minute           0
dtype: object
```

```
In [57]: # Let's compute summary statistic per a group':
flights.groupby('carrier')['dep_delay'].mean()
```

```
Out[57]: carrier
AA      8.586016
AS      5.804775
DL      9.264505
UA     12.106073
US      3.782418
Name: dep_delay, dtype: float64
```

```
In [58]: # We can use agg() methods for aggregation:
flights[['dep_delay', 'arr_delay']].agg(['min', 'mean', 'max'])
```

```
Out[58]:
```

	dep_delay	arr_delay
min	-33.000000	-75.000000
mean	9.463773	2.094537
max	1014.000000	1007.000000

```
In [59]: # An example of computing different statistics for different columns
flights.agg({'dep_delay': ['min', 'mean', 'max'], 'carrier': ['nunique']})
```

```
Out[59]:
```

	dep_delay	carrier
max	1014.000000	NaN
mean	9.463773	NaN
min	-33.000000	NaN
nunique	NaN	5.0

### 1.0.7 Basic descriptive statistics

Function	Description
min	minimum
max	maximum
mean	arithmetic mean of values
median	median
mad	mean absolute deviation
mode	mode
std	standard deviation
var	unbiased variance
sem	standard error of the mean
skew	sample skewness
kurt	kurtosis
quantile	value at %

```
In [60]: # Convenient describe() function computes a variety of statistics
         flights.dep_delay.describe()
```

```
Out[60]: count    158418.000000
         mean       9.463773
         std       36.545109
         min      -33.000000
         25%       -5.000000
         50%       -2.000000
         75%        7.000000
         max      1014.000000
         Name: dep_delay, dtype: float64
```

```
In [61]: # find the index of the maximum or minimum value
         # if there are multiple values matching idxmin() and idxmax() will return the first match
         flights['dep_delay'].idxmin() #minimum value
```

```
Out[61]: 54111
```

```
In [62]: # Count the number of records for each different value in a vector
         flights['carrier'].value_counts()
```


```
Out[62]: UA      58665
         DL      48110
         AA      32729
         US      20536
         AS        714
         Name: carrier, dtype: int64
```

### 1.0.8 Explore data using graphics


```
In [63]: #Show graphs withint Python notebook
         %matplotlib inline
```

```
In [64]: #Use matplotlib to draw a histogram of a salary data
plt.hist(df['salary'],bins=8, normed=1)
```

```
Out[64]: (array([ 7.14677085e-06,  8.73494215e-06,  1.74698843e-05,
                  8.73494215e-06,  9.52902780e-06,  6.35268520e-06,
                  3.17634260e-06,  7.94085650e-07]),
          array([ 57800.,  73945.,  90090., 106235., 122380., 138525.,
                  154670., 170815., 186960.]),
          <a list of 8 Patch objects>)
```




```
In [65]: #Use seaborn package to draw a histogram
sns.distplot(df['salary']);
```



```
In [66]: # Use regular matplotlib function to display a barplot  
df.groupby(['rank'])['salary'].count().plot(kind='bar')
```


```
Out[66]: <matplotlib.axes._subplots.AxesSubplot at 0x7ff58213f860>
```






```
In [67]: # Use seaborn package to display a barplot
sns.set_style("whitegrid")

ax = sns.barplot(x='rank', y='salary', data=df, estimator=len)
```




```
In [68]: # Split into 2 groups:  
ax = sns.barplot(x='rank', y='salary', hue='sex', data=df, estimator=len)
```




```
In [69]: #Violinplot
sns.violinplot(x = "salary", data=df)
```

```
Out[69]: <matplotlib.axes._subplots.AxesSubplot at 0x7ff5819b79e8>
```




```
In [70]: #Scatterplot in seaborn
sns.jointplot(x='service', y='salary', data=df)
```

```
Out[70]: <seaborn.axisgrid.JointGrid at 0x7ff581984550>
```




```
In [71]: #If we are interested in linear regression plot for 2 numeric variables we can use regplot
sns.regplot(x='service', y='salary', data=df)
```

```
Out[71]: <matplotlib.axes._subplots.AxesSubplot at 0x7ff58184c470>
```




```
In [72]: # box plot
sns.boxplot(x='rank',y='salary', data=df)
```

```
Out[72]: <matplotlib.axes._subplots.AxesSubplot at 0x7ff58170de80>
```




```
In [73]: # side-by-side box plot
sns.boxplot(x='rank',y='salary', data=df, hue='sex')
```

```
Out[73]: <matplotlib.axes._subplots.AxesSubplot at 0x7ff5818edc18>
```




```
In [74]: # swarm plot
sns.swarmplot(x='rank',y='salary', data=df)
```

```
Out[74]: <matplotlib.axes._subplots.AxesSubplot at 0x7ff5814f75c0>
```



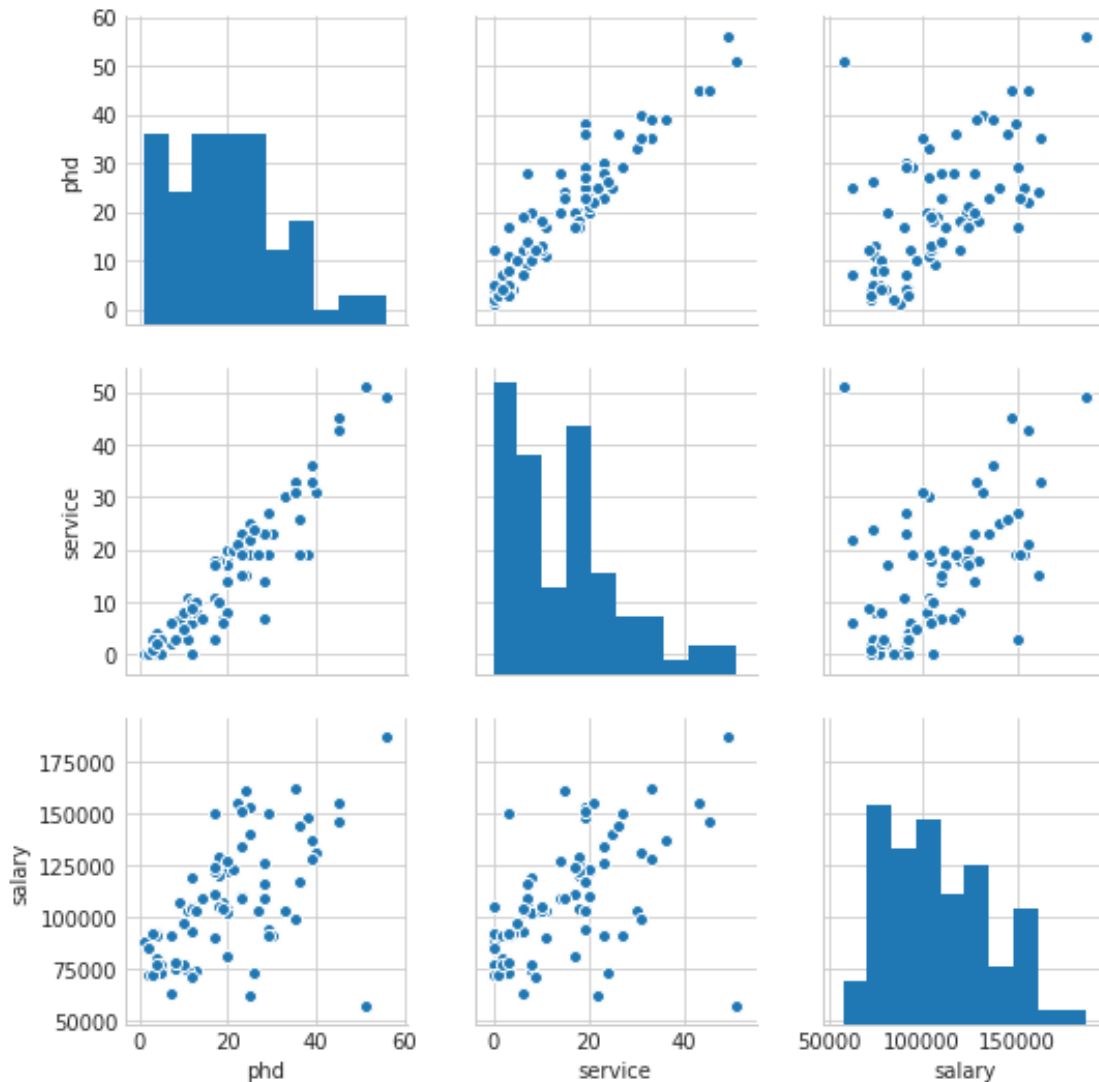
```
In [75]: #factorplot
sns.factorplot(x='carrier',y='dep_delay', data=flights, kind='bar')

Out[75]: <seaborn.axisgrid.FacetGrid at 0x7ff58178a198>
```



```
In [76]: # Pairplot
sns.pairplot(df)
```

```
Out[76]: <seaborn.axisgrid.PairGrid at 0x7ff5822296a0>
```




---

### Excercise

```
In [77]: #Using seaborn package explore the dependency of arr_delay on dep_delay (scatterplot or
sns.jointplot(x='dep_delay', y='arr_delay', data=flights)
```



Out [77]: <seaborn.axisgrid.JointGrid at 0x7ff580cb7a20>



---

## 1.1 Basic statistical Analysis

### 1.1.1 Linear Regression

```
In [78]: # Import Statsmodel functions:  
import statsmodels.formula.api as smf
```

```
In [79]: # create a fitted model  
lm = smf.ols(formula='salary ~ service', data=df).fit()
```

```
#print model summary
print(lm.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:                  salary    R-squared:                  0.283
Model:                            OLS      Adj. R-squared:              0.274
Method:                 Least Squares    F-statistic:                 30.03
Date:                Fri, 15 Sep 2017    Prob (F-statistic):          5.31e-07
Time:                  14:11:46    Log-Likelihood:              -896.72
No. Observations:                78      AIC:                        1797.
Df Residuals:                    76      BIC:                        1802.
Df Model:                        1
Covariance Type:                nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	8.935e+04	4365.651	20.468	0.000	8.07e+04	9.8e+04
service	1240.3567	226.341	5.480	0.000	789.560	1691.153

```

=====
Omnibus:                        12.741    Durbin-Watson:              1.630
Prob(Omnibus):                  0.002    Jarque-Bera (JB):           21.944
Skew:                          -0.576    Prob(JB):                   1.72e-05
Kurtosis:                      5.329    Cond. No.                   30.9
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [80]: # print the coefficients
lm.params
```

```
Out[80]: Intercept    89354.824215
service      1240.356654
dtype: float64
```

```
In [81]: #using scikit-learn:
from sklearn import linear_model
est = linear_model.LinearRegression(fit_intercept = True) # create estimator object
est.fit(df[['service']], df[['salary']])

#print result
print("Coef:", est.coef_, "\nIntercept:", est.intercept_)
```

```
Coef: [[ 1240.3566535]]
Intercept: [ 89354.82421525]
```

---

## Excercise

```
In [82]: # Build a linear model for arr_delay ~ dep_delay
lm = smf.ols(formula='arr_delay ~ dep_delay', data=flights).fit()

#print model summary
print(lm.summary())
```

OLS Regression Results

Dep. Variable:	arr_delay	R-squared:	0.794
Model:	OLS	Adj. R-squared:	0.794
Method:	Least Squares	F-statistic:	6.074e+05
Date:	Fri, 15 Sep 2017	Prob (F-statistic):	0.00
Time:	14:11:47	Log-Likelihood:	-6.8778e+05
No. Observations:	157927	AIC:	1.376e+06
Df Residuals:	157925	BIC:	1.376e+06
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	-7.4457	0.049	-152.050	0.000	-7.542	-7.350
dep_delay	1.0138	0.001	779.358	0.000	1.011	1.016

Omnibus:	38155.693	Durbin-Watson:	1.467
Prob(Omnibus):	0.000	Jarque-Bera (JB):	159178.104
Skew:	1.141	Prob(JB):	0.00
Kurtosis:	7.357	Cond. No.	38.9

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

---

### 1.1.2 Student T-test

```
In [83]: # Using scipy package:
from scipy import stats
df_w = df[ df['sex'] == 'Female']['salary']
df_m = df[ df['sex'] == 'Male']['salary']
stats.ttest_ind(df_w, df_m)
```

Out[83]: Ttest\_indResult(statistic=-2.2486865976699053, pvalue=0.027429778657910103)