

ml-ex3

August 4, 2017

0.1 Multi-class classification

This exercise is described in [ex3.pdf](#).

```
In [3]: import numpy as np
import scipy as sp
import scipy.io as sio
import matplotlib.pyplot as plt

from sklearn.linear_model import LogisticRegression, LogisticRegressionCV
from sklearn.metrics import classification_report

%matplotlib inline
```

0.1.1 Hand-written digits dataset

```
In [4]: # Load the hand-written digits dataset
digits = sio.loadmat('data/ml-ex3/ex3data1.mat')

In [5]: # Digit image data (5000 images with 400 features/pixels)
X = digits['X']

# Digit classes (1-10) where digit 0 is assigned class 10
y = digits['y'].ravel()

In [6]: # Plot three samples from each class
n_rows = 10
n_cols = 3

selected = np.linspace(0, X.shape[0] - 1, n_rows * n_cols, dtype='int16')

plt.subplots_adjust(top=.9, hspace=.4)
plt.figure(figsize=(1.8 * n_cols, 2.4 * n_rows))

for i, idx in enumerate(selected):
    plt.subplot(n_rows, n_cols, i + 1)
    plt.imshow(X[idx].reshape((20,20), order='F'), cmap=plt.cm.gray)
    plt.title(f'class={y[idx]}')
    plt.xticks(())
    plt.yticks(())
```

<matplotlib.figure.Figure at 0x107172ba8>

class=10



class=10



class=10



class=1



class=1



class=1



class=2



class=2



class=2



class=3



class=3



class=3



class=4



class=4



class=4



class=5



class=5



class=5



class=6



class=6



class=6



class=7



class=7



class=7



class=8



class=8



class=8



class=9



class=9



class=9



0.1.2 Logistic regression

```
In [7]: # Train a logistic regression classifier with C=1.0 (inverse regularization strength)
        clf = LogisticRegression(C=1.0, solver='newton-cg')
        clf.fit(X, y)
```

```
Out[7]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                           intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
                           penalty='l2', random_state=None, solver='newton-cg', tol=0.0001,
                           verbose=0, warm_start=False)
```

```
In [8]: # Classification accuracy on training set
        clf.score(X, y)
```

```
Out[8]: 0.9446
```

```
In [9]: # Train a logistic regression classifier, running
        # built-in cross validation for selecting the best
        # C value
        clf_cv = LogisticRegressionCV(Cs=[1e-1, 1e0, 1e1], solver='newton-cg')
        clf_cv.fit(X, y)
```

```
Out[9]: LogisticRegressionCV(Cs=[0.1, 1.0, 10.0], class_weight=None, cv=None,
                              dual=False, fit_intercept=True, intercept_scaling=1.0,
                              max_iter=100, multi_class='ovr', n_jobs=1, penalty='l2',
                              random_state=None, refit=True, scoring=None, solver='newton-cg',
                              tol=0.0001, verbose=0)
```

```
In [10]: # Classification accuracy on training set
         # using the model that was selected during
         # cross-validation
         clf_cv.score(X, y)
```

```
Out[10]: 0.94079999999999997
```

0.1.3 Neural networks

```
In [11]: # Load the parameters of a pre-trained neural network
        weights = sio.loadmat('data/ml-ex3/ex3weights.mat')
```

```
        Theta1 = weights['Theta1']
        Theta2 = weights['Theta2']
```

```
        Theta1.shape, Theta2.shape
```

```
Out[11]: ((25, 401), (10, 26))
```

```

In [12]: # Number of samples (first dimension of design matrix X)
         n_samples = X.shape[0]

         # sigmoid function
         sigmoid = sp.special.expit

In [13]: # Input layer data
         A1 = np.c_[np.ones(n_samples), X]

In [14]: # Hidden layer activations
         Z2 = A1.dot(Theta1.T)
         A2 = np.c_[np.ones(n_samples), sigmoid(Z2)]

In [15]: # Output layer activations
         Z3 = A2.dot(Theta2.T)
         A3 = sigmoid(Z3)

In [16]: # Predicted class is index of highest value per row + 1
         y_pred = np.argmax(A3, axis=1) + 1

In [17]: # Print classification report showing the main classification metrics
         print(classification_report(y, y_pred, target_names=['1', '2', '3', '4', '5', '6', '7', '8', '9']

```

	precision	recall	f1-score	support
1	0.9684	0.9820	0.9752	500
2	0.9818	0.9700	0.9759	500
3	0.9776	0.9600	0.9687	500
4	0.9699	0.9680	0.9690	500
5	0.9723	0.9840	0.9781	500
6	0.9782	0.9860	0.9821	500
7	0.9778	0.9700	0.9739	500
8	0.9781	0.9820	0.9800	500
9	0.9657	0.9580	0.9618	500
10	0.9822	0.9920	0.9871	500
avg / total	0.9752	0.9752	0.9752	5000