

# ml-ex4

August 4, 2017

## 0.1 Neural networks learning

This exercise is described in [ex4.pdf](#).

```
In [1]: import scipy.io as sio
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.neural_network import MLPClassifier
from sklearn.model_selection import StratifiedShuffleSplit, GridSearchCV

%matplotlib inline
```

```
In [2]: # Load the hand-written digits dataset
digits = sio.loadmat('data/ml-ex4/ex4data1.mat')
```

```
In [3]: # Digit image data (5000 images with 400 features/pixels)
X = digits['X']

# Digit classes (1-10) where digit 0 is assigned class 10
y = digits['y'].ravel()
```

Sample images are shown in [Exercise 3 notebook](#).

```
In [4]: # Create a feed-forward neural network with a single hidden layer
# with 25 logistic units. Input layer size is derived from number
# of features in X. Output layer size is derived from number of
# classes in y.
clf = MLPClassifier((25,), activation='logistic', solver='lbfgs')
clf.fit(X, y)
```

```
Out[4]: MLPClassifier(activation='logistic', alpha=0.0001, batch_size='auto',
beta_1=0.9, beta_2=0.999, early_stopping=False, epsilon=1e-08,
hidden_layer_sizes=(25,), learning_rate='constant',
learning_rate_init=0.001, max_iter=200, momentum=0.9,
nesterovs_momentum=True, power_t=0.5, random_state=None,
shuffle=True, solver='lbfgs', tol=0.0001, validation_fraction=0.1,
verbose=False, warm_start=False)
```

```
In [5]: # Classification accuracy on training set
        clf.score(X, y)
```

```
Out[5]: 1.0
```

```
In [12]: # For cross validation, use two stratified randomized folds
         # where the test set size = 0.1 * dataset size. Statified
         # means that the folds are made by preserving the percentage
         # of samples for each class.
         cv = StratifiedShuffleSplit(2, test_size=0.1, random_state=0)

         # Run a grid search to find the best value values for regularization
         # parameter alpha using the cross validator (cv) defined above.
         gs = GridSearchCV(clf, param_grid={'alpha':[1e-2, 1e-1, 1e0, 1e1]}, cv=cv)
         gs.fit(X, y)
```

```
Out[12]: GridSearchCV(cv=StratifiedShuffleSplit(n_splits=2, random_state=0, test_size=0.1,
        train_size=None),
        error_score='raise',
        estimator=MLPClassifier(activation='logistic', alpha=0.0001, batch_size='auto',
        beta_1=0.9, beta_2=0.999, early_stopping=False, epsilon=1e-08,
        hidden_layer_sizes=(25,), learning_rate='constant',
        learning_rate_init=0.001, max_iter=200, momentum=0.9,
        nesterovs_momentum=True, power_t=0.5, random_state=None,
        shuffle=True, solver='lbfgs', tol=0.0001, validation_fraction=0.1,
        verbose=False, warm_start=False),
        fit_params={}, iid=True, n_jobs=1,
        param_grid={'alpha': [0.01, 0.1, 1.0, 10.0]},
        pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
        scoring=None, verbose=0)
```

```
In [7]: # Show the cross validation results in a pandas data frame
         # - column mean_train_score: mean classification accuracy on training set
         # - column mean_test_score: mean classification accuracy on test set
         pd.DataFrame(gs.cv_results_)
```

```
Out[7]:
```

|   | mean_fit_time | mean_score_time | mean_test_score | mean_train_score | \ |
|---|---------------|-----------------|-----------------|------------------|---|
| 0 | 1.437032      | 0.001812        | 0.919           | 1.000000         |   |
| 1 | 2.473062      | 0.001954        | 0.937           | 1.000000         |   |
| 2 | 4.012578      | 0.002556        | 0.946           | 0.997556         |   |
| 3 | 4.014157      | 0.002660        | 0.920           | 0.943222         |   |

  

|   | param_alpha | params          | rank_test_score | split0_test_score | \ |
|---|-------------|-----------------|-----------------|-------------------|---|
| 0 | 0.01        | {'alpha': 0.01} | 4               | 0.910             |   |
| 1 | 0.1         | {'alpha': 0.1}  | 2               | 0.924             |   |
| 2 | 1           | {'alpha': 1.0}  | 1               | 0.936             |   |
| 3 | 10          | {'alpha': 10.0} | 3               | 0.912             |   |

  

|  | split0_train_score | split1_test_score | split1_train_score | std_fit_time | \ |
|--|--------------------|-------------------|--------------------|--------------|---|
|--|--------------------|-------------------|--------------------|--------------|---|

|   |          |       |          |          |
|---|----------|-------|----------|----------|
| 0 | 1.000000 | 0.928 | 1.000000 | 0.004382 |
| 1 | 1.000000 | 0.950 | 1.000000 | 0.046205 |
| 2 | 0.997333 | 0.956 | 0.997778 | 0.067075 |
| 3 | 0.942889 | 0.928 | 0.943556 | 0.046254 |

|   | std_score_time | std_test_score | std_train_score |
|---|----------------|----------------|-----------------|
| 0 | 0.000024       | 0.009          | 0.000000        |
| 1 | 0.000198       | 0.013          | 0.000000        |
| 2 | 0.000335       | 0.010          | 0.000222        |
| 3 | 0.000469       | 0.008          | 0.000333        |

```
In [8]: # Best classifier is the neural network with alpha=1.0
```

```
clf_best = gs.best_estimator_  
clf_best
```

```
Out[8]: MLPClassifier(activation='logistic', alpha=1.0, batch_size='auto', beta_1=0.9,  
                      beta_2=0.999, early_stopping=False, epsilon=1e-08,  
                      hidden_layer_sizes=(25,), learning_rate='constant',  
                      learning_rate_init=0.001, max_iter=200, momentum=0.9,  
                      nesterovs_momentum=True, power_t=0.5, random_state=None,  
                      shuffle=True, solver='lbfgs', tol=0.0001, validation_fraction=0.1,  
                      verbose=False, warm_start=False)
```

```
In [9]: # Classification accuracy on training set
```

```
clf_best.score(X, y)
```

```
Out[9]: 0.99739999999999995
```

```
In [10]: # Obtain weight matrix between input and hidden layer
```

```
# (input layer size = 400, hidden layer size = 25)
```

```
Theta1 = clf_best.coefs_[0].T
```

```
Theta1.shape
```

```
Out[10]: (25, 400)
```

```
In [11]: # Visualize input weights into hidden units (400 per hidden unit)
```

```
n_rows = 5
```

```
n_cols = 5
```

```
plt.subplots_adjust(top=.9, hspace=.4)
```

```
plt.figure(figsize=(1.8 * n_cols, 2.4 * n_rows))
```

```
for i, row in enumerate(Theta1):
```

```
    plt.subplot(n_rows, n_cols, i + 1)
```

```
    plt.imshow(row.reshape((20,20), order='F'), cmap=plt.cm.gray)
```

```
    plt.title(f'hidden unit {i + 1}')
```

```
    plt.xticks(())
```

```
    plt.yticks(())
```

<matplotlib.figure.Figure at 0x10684c860>

