

# ml-ex5

August 4, 2017

## 0.1 Bias and variance

This exercise is described in [ex5.pdf](#).

```
In [1]: import numpy as np
import scipy.io as sio
import matplotlib.pyplot as plt

from sklearn.linear_model import Ridge
from sklearn.model_selection import learning_curve, train_test_split, ShuffleSplit
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import PolynomialFeatures, StandardScaler

%matplotlib inline
```

### 0.1.1 Prepare dataset

```
In [2]: # Load pre-split dataset provided by exercise
data = sio.loadmat('data/ml-ex5/ex5data1.mat')

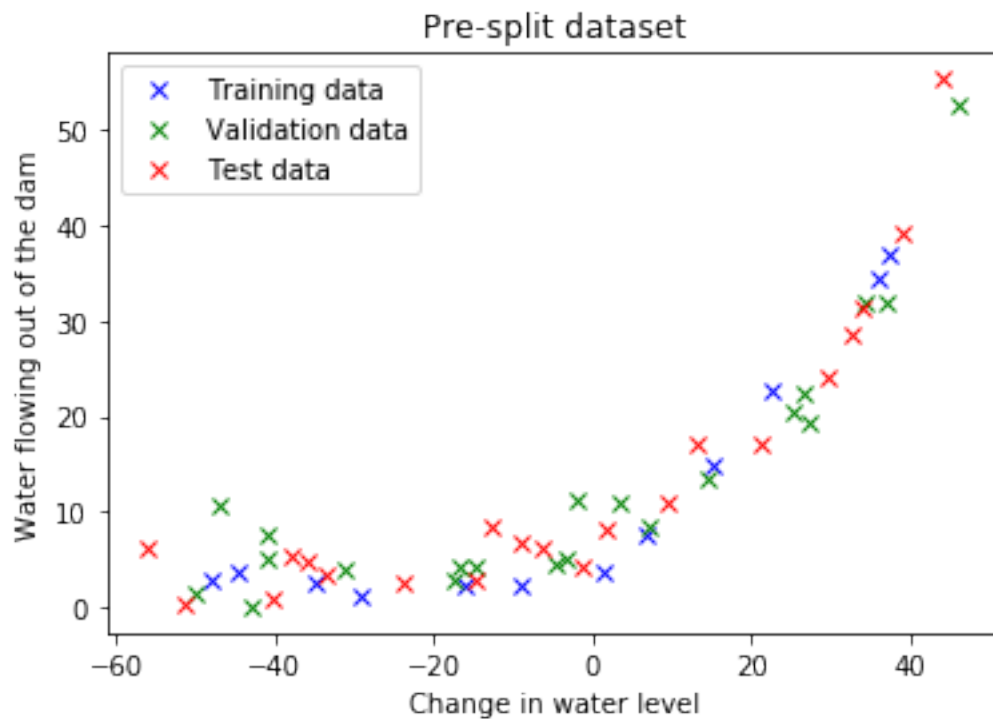
In [3]: # Training set
X_train_0 = data['X']
y_train_0 = data['y']

# Validation set
X_cval_0 = data['Xval']
y_cval_0 = data['yval']

# Test set
X_test_0 = data['Xtest']
y_test_0 = data['ytest']

In [4]: # Plot the three pre-split datasets
plt.plot(X_train_0, y_train_0, 'bx', label='Training data')
plt.plot(X_cval_0, y_cval_0, 'gx', label='Validation data')
plt.plot(X_test_0, y_test_0, 'rx', label='Test data')
plt.xlabel('Change in water level')
plt.ylabel('Water flowing out of the dam')
plt.title('Pre-split dataset')
plt.legend()
```

Out [4]: <matplotlib.legend.Legend at 0x109875f60>



```
In [5]: # Number of training data
X_train_0.shape[0]
```

Out [5]: 12

```
In [6]: # Number of validation data
X_cval_0.shape[0]
```

Out [6]: 21

```
In [7]: # Number of test data
X_test_0.shape[0]
```

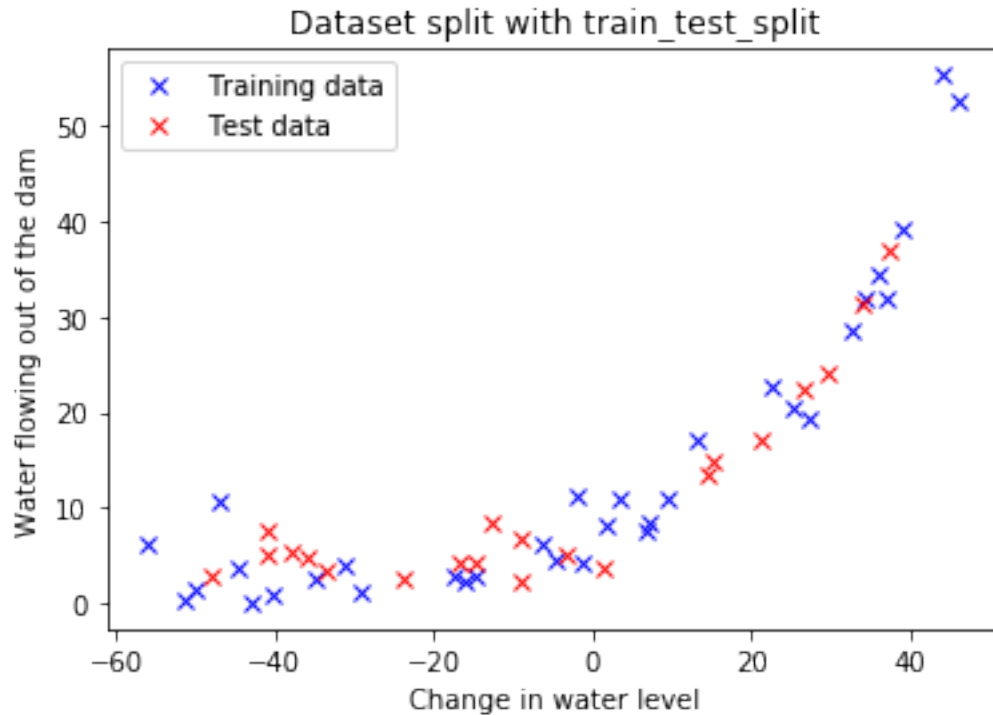
Out [7]: 21

```
In [8]: # Concatenate pre-split datasets into a single dataset
# (use scikit-learn utilities later to split datasets).
X = np.concatenate([X_train_0, X_cval_0, X_test_0])
y = np.concatenate([y_train_0, y_cval_0, y_test_0]).ravel()
```

```
In [9]: # Randomly split into training and test set. Test set size = 21.
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=21, random_state=42)
```

```
In [10]: # Plot the datasets obtained with train_test_split
plt.plot(X_train, y_train, 'bx', label='Training data')
plt.plot(X_test, y_test, 'rx', label='Test data')
plt.xlabel('Change in water level')
plt.ylabel('Water flowing out of the dam')
plt.title('Dataset split with train_test_split')
plt.legend()
```

```
Out[10]: <matplotlib.legend.Legend at 0x109a564e0>
```



## 0.1.2 Utility functions

```
In [11]: # Plot training data and predictions by a trained model
def plot_training_data_and_predictions(X_train, y_train, X_pred, y_pred, title=''):
    plt.plot(X_train, y_train, 'bx', label='Training data')
    plt.plot(X_pred, y_pred, 'b--', label='Prediction')
    plt.xlabel('Change in water level')
    plt.ylabel('Water flowing out of the dam')
    plt.title(title)
    plt.legend()

# Plot learning curves obtained by training on different training set sizes
def plot_learning_curves(train_sizes, train_scores, test_scores, y_min=-1, y_max=1, title=''):
    plt.plot(train_sizes, train_scores, 'b-', label='Training score')
```

```

plt.plot(train_sizes, test_scores, 'g-', label='Test score (CV)')

plt.grid(True, axis='y')
plt.xlabel('Training set size')
plt.ylabel('Score')
plt.xlim(xmin=1)
plt.ylim(y_min, y_max)
plt.title(title)
plt.legend()

# Generate learning curves by training regressor on different training set sizes
def compute_learning_curves(regressor, X_train, y_train, test_size):
    # Train/test split using test_size examples
    cv = ShuffleSplit(n_splits=50, test_size=test_size, random_state=0)
    # Relative training set sizes
    train_sizes_rel = np.linspace(.1, 1.0, 10)
    # Generate learning curves
    train_sizes, train_scores, test_scores = learning_curve(regressor, X_train, y_train)
    # Return training set sizes and average learning values
    return [train_sizes, np.mean(train_scores, axis=1), np.mean(test_scores, axis=1)]

```

### 0.1.3 Linear Regression

```

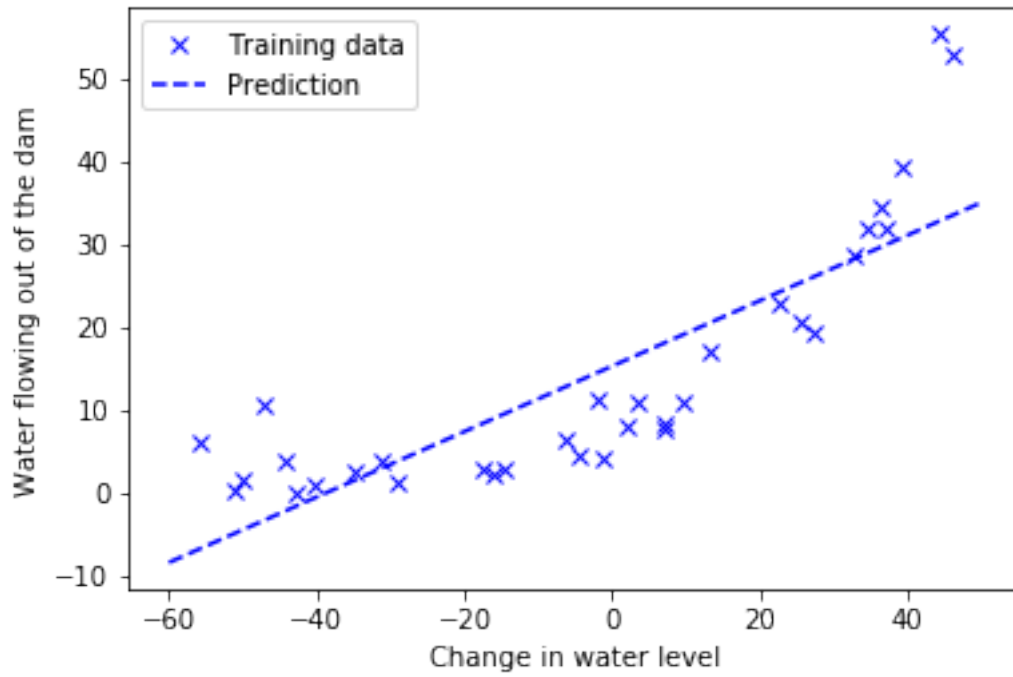
In [12]: # Linear regression with default regularization strength (alpha=1.0)
         alpha = 1.0

         model = Ridge(alpha=alpha)
         model.fit(X_train, y_train)

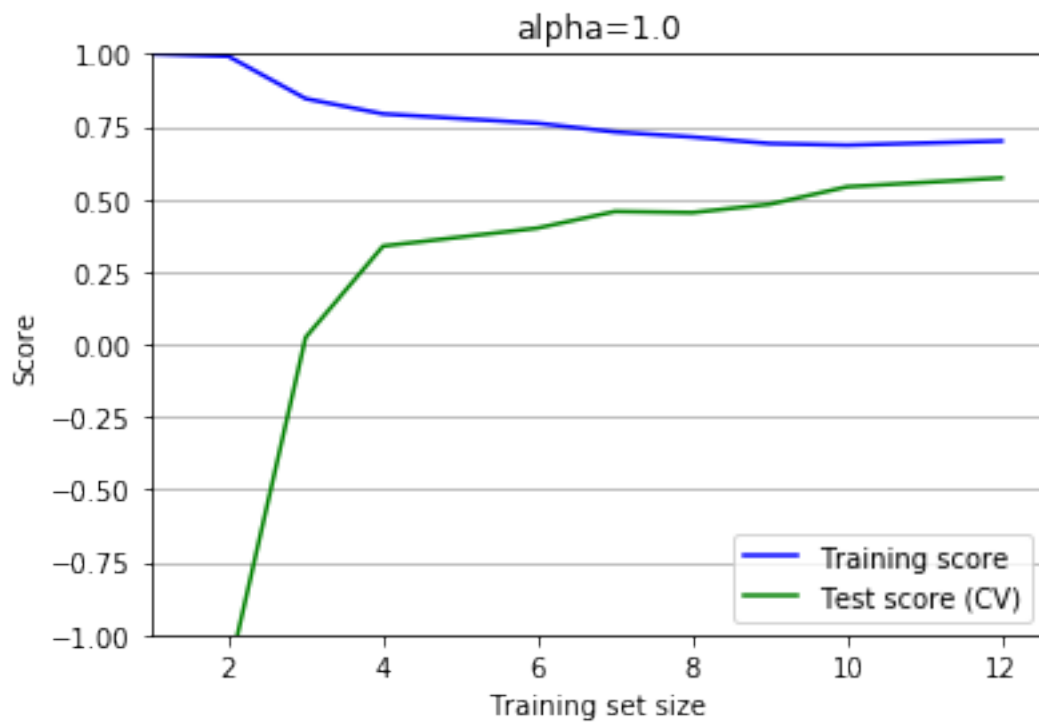
         X_pred = np.array([[ -60], [ 50]])
         y_pred = model.predict(X_pred)

         plot_training_data_and_predictions(X_train, y_train, X_pred, y_pred)

```



```
In [13]: train_sizes, train_scores, test_scores = compute_learning_curves(model, X_train, y_train)
         plot_learning_curves(train_sizes, train_scores, test_scores, title=f'alpha={alpha}')
```



```
In [14]: # Score on test set
         model.score(X_test, y_test)
```

```
Out[14]: 0.60084127499786522
```

### 0.1.4 Polynomial regression

```
In [15]: # Creates a linear regressor with given regularization strength alpha.
         # Ploynomial features of degree 8 are added and scaled before running
         # regulaized linear regression.
         def regressor(alpha):
             return Pipeline(steps=[
                 ('poly', PolynomialFeatures(degree=8)),
                 ('scaler', StandardScaler()),
                 ('lreg', Ridge(alpha=alpha))
             ])
```


### Overfit example

```
In [16]: alpha = 0.001


         model = regressor(alpha=alpha)
         model.fit(X_train, y_train)

         X_pred = np.linspace(-70, 60, 100).reshape(-1,1)
         y_pred = model.predict(X_pred)

         plot_training_data_and_predictions(X_train, y_train, X_pred, y_pred, f'alpha={alpha}')
```



```
In [17]: train_sizes, train_scores, test_scores = compute_learning_curves(model, X_train, y_train)
         plot_learning_curves(train_sizes, train_scores, test_scores, y_min=-15, title=f'alpha={alpha}')
         plt.show()
```



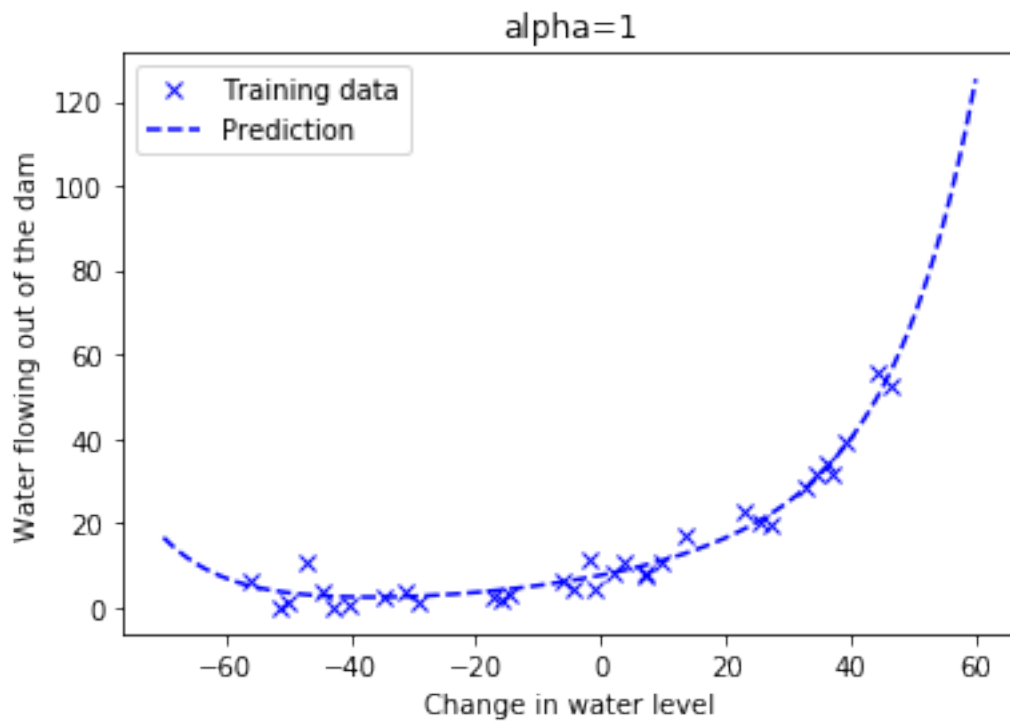
## Good fit example

```
In [18]: alpha = 1
```

```
model = regressor(alpha=alpha)
model.fit(X_train, y_train)
```

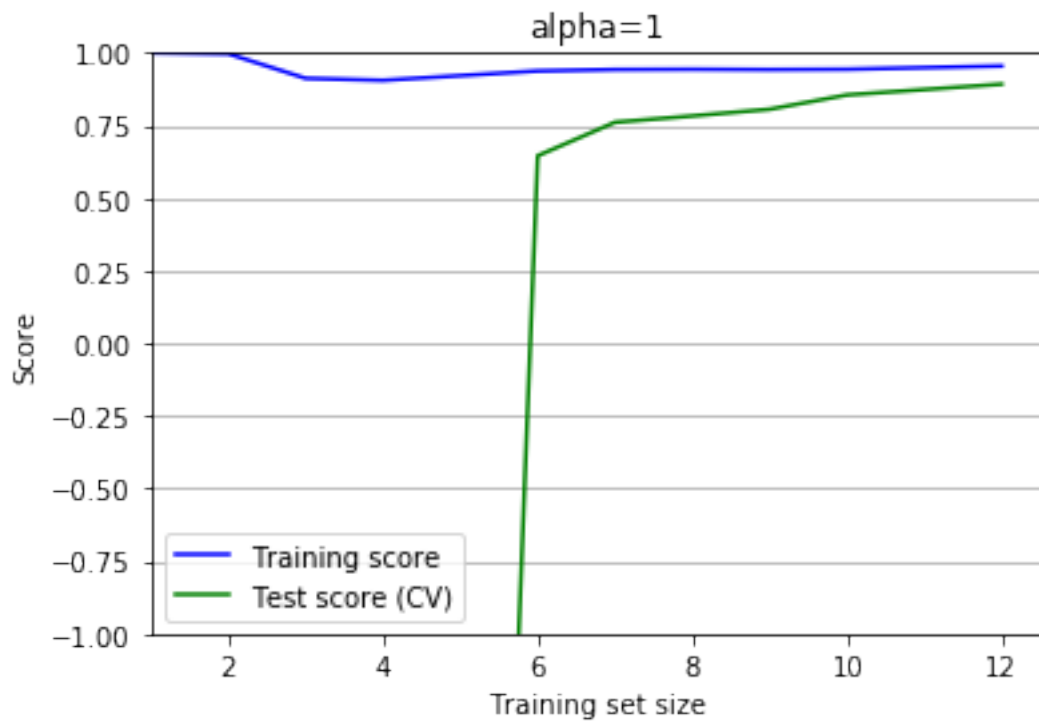
```
X_pred = np.linspace(-70, 60, 100).reshape(-1,1)
y_pred = model.predict(X_pred)
```

```
plot_training_data_and_predictions(X_train, y_train, X_pred, y_pred, f'alpha={alpha}')
```



```
In [19]: train_sizes, train_scores, test_scores = compute_learning_curves(model, X_train, y_train)
plot_learning_curves(train_sizes, train_scores, test_scores, title=f'alpha={alpha}')
```





```
In [20]: # Score on test set
         model.score(X_test, y_test)
```

```
Out[20]: 0.95207578600535314
```

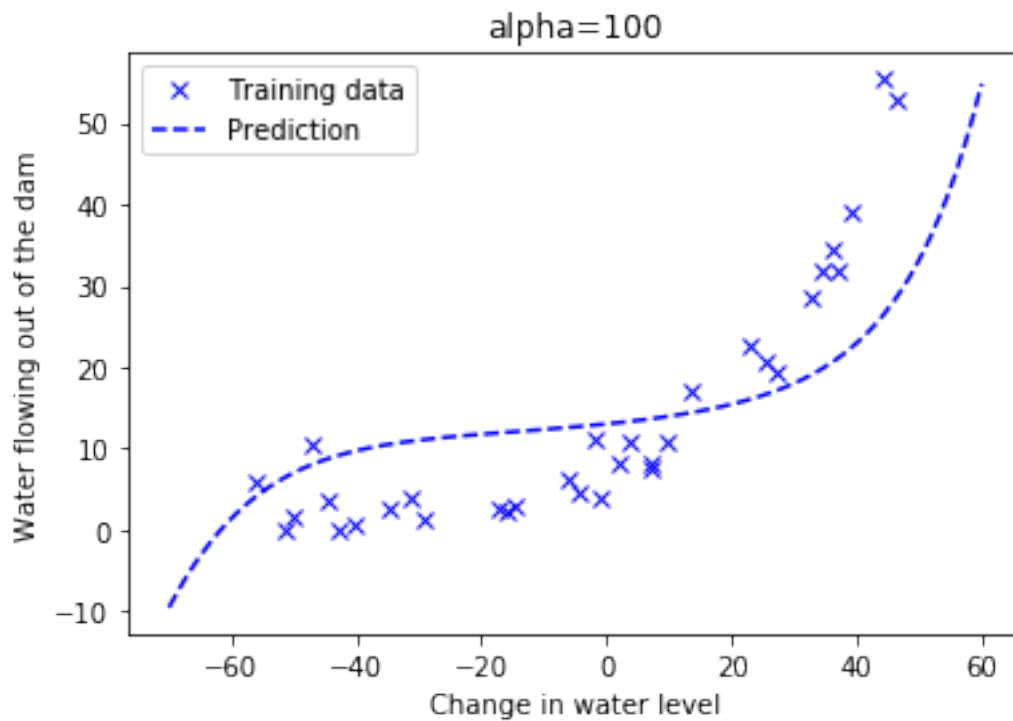
### Underfit example

```
In [21]: alpha = 100
```


```
model = regressor(alpha=alpha)
model.fit(X_train, y_train)
```

```
X_pred = np.linspace(-70, 60, 100).reshape(-1,1)
y_pred = model.predict(X_pred)
```

```
plot_training_data_and_predictions(X_train, y_train, X_pred, y_pred, f'alpha={alpha}')
```



```
In [22]: train_sizes, train_scores, test_scores = compute_learning_curves(model, X_train, y_train)
         plot_learning_curves(train_sizes, train_scores, test_scores, title=f'alpha={alpha}')
```



```
In [23]: # Score on test set  
         model.score(X_test, y_test)
```

```
Out[23]: 0.45166186769096572
```